

Part V

Memory Management

Chapter 9

Memory

9.1 PROC File System

Give the following command without any arguments and press Enter:

```
mount
```

You will see a list of all partitions that are currently mounted on your system. If you look closely, there will be an entry called PROC. PROC is a special file system and is an open door to the Linux kernel. You can peep in here to see what the kernel is doing. The proc file system can be accessed by changing into the /proc directory.

```
cd /proc
```

Here, you will apparently see plenty of files. But these are not files, rather they are parameters, data structures, and statistics collected from the kernel and which appear like files inside the /proc directory. The contents of this directory are created 'on-the-fly' by the kernel whenever we read its contents. The existence of "files" in this directory only appear the instance you request it. Otherwise, the directory is ***EMPTY***.

9.1.1 Exploring PROC

So let's start exploring this file system. Type the following command:

```
ls -l /proc/version
```

Note the following about above command:

Q1 What is the file size?

Q2 What is the modification time?

Q3 Now check the current time on your watch?

Do you notice that the modification time is the same as the time you entered the command. If not, wait for 1 minute and run the command again. You will notice that the modification time has changed yet you haven't even touched the file so far. This proves that the information appears on the fly and that nothing exists in the directory as such.

Now let's view the contents of this file. Type the following:

```
cat /proc/version
```

You can see version numbers such as the kernel, and the compiler which compiled the kernel. Give the ls command again. Note the following again:

Q What is the file size? What is the modification time? How come so much information was just showed as "contents" of the file yet the file size is zero? Again, this is confirmation of previous point.

9.1.2 Processes in PROC

The numbers you see inside the PROC file system are directory representations for each process currently running on your system. These numbers are the process id's. If you give the ls command repeatedly, you will notice that these directory representations are changing dynamically. Choose a process ID and enter that directory using:

```
cd /proc/<number>
```

View the contents. You should be able to see files like `cmdline`, `cwd`, `environ`, `exe`, `fd`, `statm`, `status`, etc. Of these, `fd` contains the number of open file descriptors for any given process.

Using the `ps` command, find out the process id of your current shell.

```
ps
```

Now give the following command:

```
ls -l /proc/<your shell process id>/fd
```

Q What open file descriptors can you see from the output?

Q What do the arrow operator (`->`) represent?

Open a new terminal window and type the following:

```
echo "Hello World" >> /proc/
    <your shell process id>/fd/1
```

Go back to the previous terminal window.

Q What just happened? We had two terminal windows open. Each terminal window was one instance of shell process (so 2 processes). We gave a command in one shell and the output appeared in the other shell. Isn't this communication between two processes? The IPC method of communication was through a "file".

9.1.3 Process Memory in PROC

Look at the output of the `statm` file. It should display a list of seven numbers separated by spaces. Each number is the number of page frames of memory used by a process in a particular category. The categories are:

- Total process size
- Size of process currently resident in main memory
- Size of memory shared with other processes
- Text section size of a process

- Size of shared libraries
- Memory used by the process for its stack
- Number of dirty pages (pages of memory that have been modified by the program)

Check these information for your currently logged in shell.

```
cat /proc/<your shell process id>/statm
cat /proc/<your shell process id>/status
```

9.1.4 Overall Memory Usage in PROC

Overall memory usage can be seen from

```
cat /proc/meminfo
```

Q Find out the total memory of your machine

Q How much of it is free

Q How much is your swap space

Q How much of that swap space is free from

```
meminfo
```

9.2 Exploring Memory

9.2.1 Logical and Physical Addresses

Write the following simple C program.

```
#include <unistd.h>
#include <stdio.h>
#include <fcntl.h>
int main()
{
    int fd, bw, br;
    char *buffer=(char*)calloc(NULL, BUFSIZ);
    fd = open("/tmp/foo.txt",O_CREAT|O_RDWR);
    br = read(0,buffer, BUFSIZ);
    bw = write(fd,buffer,BUFSIZ);
    close(fd);
    return 0;
}
```

You can see usage of I/O system calls such as `open()`, `read()`, `write()` and `close()`. You can also see a system call called `calloc()`. `Calloc()` is allocating un-used space for an array of n number of elements. We will see description of `calloc()` and its sister system calls shortly.

9.2.1.1 Logical Addresses

Compile this code using:

```
gcc code.c -c
```

and then

```
gcc code.c -o code -g
```

The `-g` is for adding debug symbols.

Q What is the size of the executable file with and without the `-g` argument?

We are now going to use the `objdump` utility to disassemble (`-d`) the program and view the information about the object file.

Run the following:

```
objdump -D code.o | less
```

Note: If the output is too much, you can enter

```
objdump -D code.o >
/location/of/directory/code.txt
```

You can then view the contents of the output by opening `code.txt` in any text editor.

Find out the following (You will have to learn to scroll the huge list of addresses up and down and learn to catch addresses):

Q The address on which your first instruction of your program starts. Note it down.

Q The address at which the `main()` function starts. Note this down as well.

Q Identify the instructions for `calloc()`, `open()`, `read()`, `write()` and `close()` calls and note down the addresses on which these instructions appear. This should be in the main function.

Q What do you think? Are the addresses to the left logical or physical addresses?

Q Find out where and in what section the `write()` call is referring to? Go to that section. How many lines of code are in that section (making a rough estimate)?

Q Do you think these lines of code are the actual implementation of the `write` call? `Write()` call is from the `unistd.h` library. So do you think we are currently looking at the function from the `unistd.h` library?

9.2.1.2 Physical Addresses

Now exit and recompile your code with the static option.

```
gcc code.c -o code -static
```

Q What is the size of the executable file with and without the `-g` argument?

Use `Objdump` again.

```
objdump -D code | less
```

Try to answer the following and compare with the information you noted down earlier.

Q What is the starting point of your first instruction again?

Q What is the address of the main function?

Q Do you think this time it is a logical or a physical address?

Q At what instruction addresses are the different calls located (`calloc`, `open`, `read`, `write`, `close`, etc)

Q In which section is the `write()` call referring to? Where is it referring to in memory address? Go to that address !!!

Q Do you think this is the actual implementation area of the `write()` call?

9.3 System Calls for Memory Handling**9.3.1 Introduction**

For sake of definition, a process is a running program. It means an OS has loaded a program into memory, has arranged some environment for it, and has started running it. Memory is required for the following areas in a process:

- Text Section: The portion of process where the actual executable code resides.
- Data Section: Containing global data that are to be generated during run-time when the program starts.

- **Heap:** The region from where memory can be dynamically allocated to a process.
- **Stack Section:** Where local variables (non-static) and function calls are implemented.

To have a general overview of how much size is allocated to a process, do the following simple test.

1. Create a simple Hello World
2. Compile it into an executable called Hello
3. Then, type the command: `ls -l Hello` and note down the size of the executable
4. Then, type the command: `size Hello` and note down the total size of all sections

```
ls -l Hello
size Hello
```

You will notice that a lot of the size given in the `ls` command is nowhere to be seen for the `size` command. Where does it go? Look at option 3: Heap.

9.3.2 Memory Allocation

We have some library functions in C that are used for allocating additional memory to a process. These allocations are done at run-time. Dynamic memory allocation are performed by the calls `malloc()` or `calloc()`. When successful, these will return pointers to the portion of memory just allocated. Once a portion of memory is allocated, we can change its size dynamically using the `realloc()` call. Finally, memory released is freed up using the `free()` call.

9.3.2.1 Malloc()

Initially memory is allocated using `malloc`. The value passed is the total number of bytes requested. If memory is not allocated, `NULL` is returned. If memory is allocated, a pointer to 1st address is returned. Try multiplying `sizeof` by a large number to see if it can allocate that much memory.

```
#include <stdio.h>
#include <stdlib.h>
int main()
{
    struct coord
    {
        int x, y, z;
    };
    struct coord *p;
    p = malloc(sizeof(struct coord));
    if (p==NULL)
    {
        printf("Failed");
    }
    else
    {
        printf("\n%d bytes allocated at address %d\n",
            sizeof(struct coord), p);
    }
    return 0;
}
```

9.3.2.2 Free()

Once we are done with memory, we can return it back using the `free()` function.

Continuing with previous code, we add the following:

```
free(p);
p = NULL;
```

The first line will free the memory address returned to pointer `p`. This means it is marked for use if additional memory is required for the process. However, variables still pointing to it can still use it until the time it is overwritten. This is known as a dangling pointer. It is good idea to set the pointers to `NULL` along with usage of `free`.

Try free two times and see the output.

9.3.2.3 Realloc()

Using `realloc()` call, we can dynamically change the size of the block of memory allocated earlier on. Its usage is simple:

```
p = realloc(p, sizeof(struct coord)*10);
```

This will take the current block returned to pointer `p`, and change the size to whatever we specify as the second parameter. There is, however, a problem with the above statement .. can you identify it?

9.3.2.4 Calloc()

When we allocate space using `malloc()`, it is not initialized. `Calloc()` call is just a wrapper around `malloc()`. Not only does it allocate the required space, it also initializes it to 0.

```
p = calloc(1, sizeof(struct coord));
```

Where, the first parameter is the number of members, and the second parameter is the size required for a member.

9.3.3 Exercise

Study the code below. What is wrong with it?

```
#include <stdio.h>
#include <stdlib.h>
void f(void)
{
    void* s;
    s = malloc(50);
    return;
}
int main(void)
{
    while (1)
        f();
    return 0;
}
```